Deep Learning for Computer Vision

Self-Attention and Transformers

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad

भारतीय प्रौद्योगिकी संस्थान हैदराबाद
Indian Institute of Technology Hyderabad

Review: Question

Other ways to evaluate Visual Dialog systems?

Review: Question

Other ways to evaluate Visual Dialog systems?

Look to NLP for consensus metrics that measure consensus between answers generated by model and a set of relevant answers; see Massiceti et al, A Revised Generative Evaluation of Visual Dialogue, arXiv 2020

Acknowledgements

- Most of this lecture's slides are based on Jay Alammar's article on "The Illustrated Transformer"
- Unless explicitly specified, assume that content and figures are either directly taken or adapted from above source

Motivation for Transformers

- Sequential computation prevents parallelization

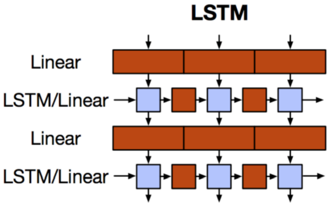

Motivation for Transformers

- Sequential computation prevents parallelization

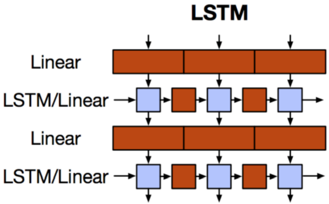

- Despite GRUs and LSTMs, RNNs still need attention mechanism to deal with long-range dependencies – path length for co-dependent computation between states grows with sequence length

Motivation for Transformers

- Sequential computation prevents parallelization

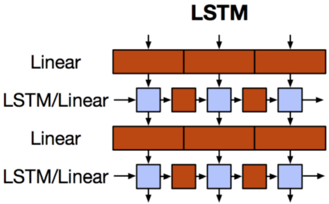

- Despite GRUs and LSTMs, RNNs still need attention mechanism to deal with long-range dependencies – path length for co-dependent computation between states grows with sequence length
- But if attention gives us access to any state, maybe we don't need the RNN?!

Credits: Richard Socher (Stanford CS224n)

Transformers[1]

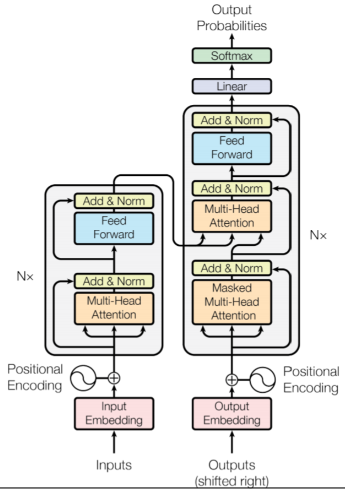

- The work "Attention is All you Need" (Vaswani et al, NeurIPS 2017) first made it possible to do Seq2Seq modeling without RNNs

[1] Vaswani et al, Attention is All You Need, NeurIPS 2017

Transformers[1]

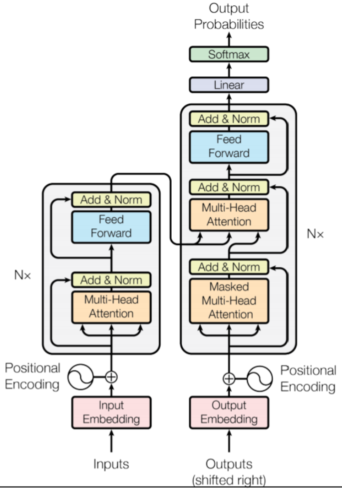

- The work "Attention is All you Need" (Vaswani et al, NeurIPS 2017) first made it possible to do Seq2Seq modeling without RNNs
- Proposed **transformer model**, entirely built on **self-attention mechanism** without using sequence-aligned recurrent architectures

[1] Vaswani et al, Attention is All You Need, NeurIPS 2017

Transformers[1]

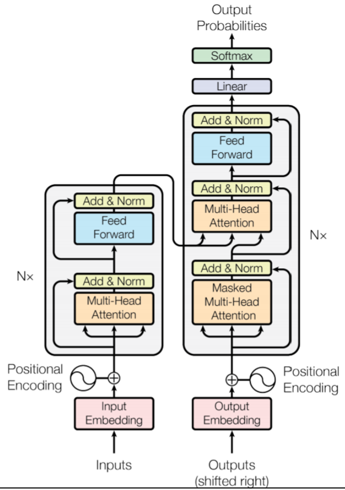

- The work "Attention is All you Need" (Vaswani et al, NeurIPS 2017) first made it possible to do Seq2Seq modeling without RNNs
- Proposed **transformer model**, entirely built on **self-attention mechanism** without using sequence-aligned recurrent architectures
- Key components:
  - Self-Attention
  - Multi-Head Attention
  - Positional Encoding
  - Encoder-Decoder Architecture

[1] Vaswani et al, Attention is All You Need, NeurIPS 2017

Transformers in a Nutshell

Transformers in a Nutshell

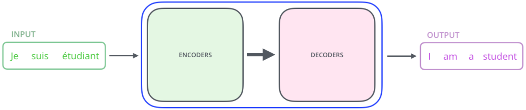

Transformers in a Nutshell

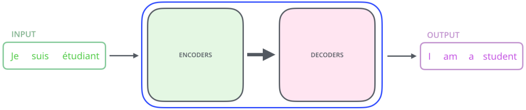

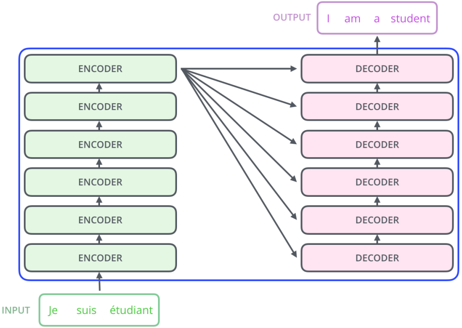

Transformers in a Nutshell

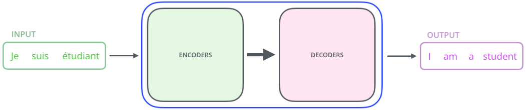

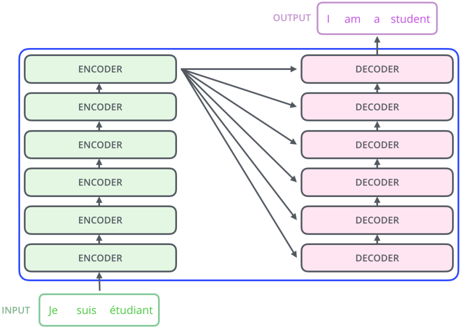

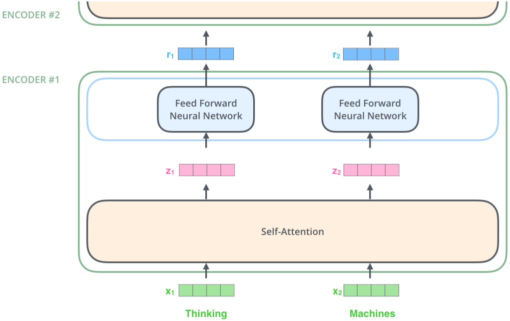

Self-Attention

- Consider two input sentences we want to translate:

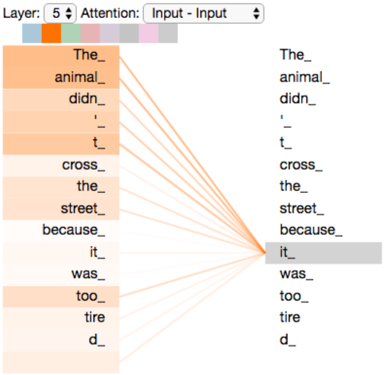

Self-Attention

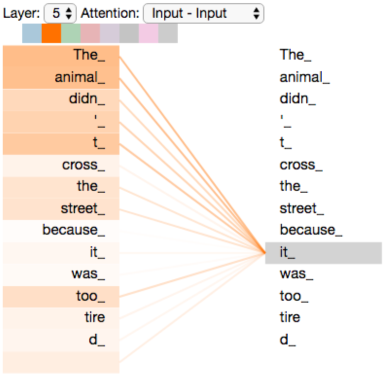

- Consider two input sentences we want to translate:
  - *The animal didn't cross the street because it was too tired*

Self-Attention

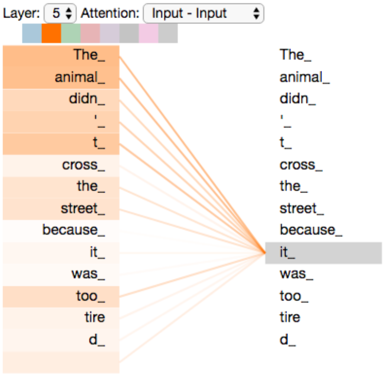

- Consider two input sentences we want to translate:
  - *The animal didn't cross the street because it was too tired*
  - *The animal didn't cross the street because it was too wide*

Self-Attention

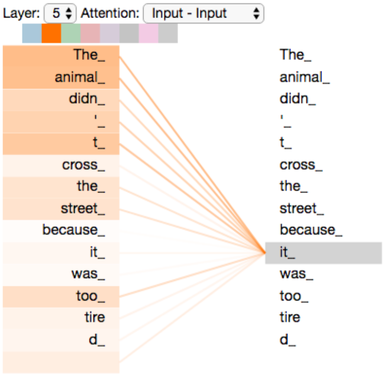

- Consider two input sentences we want to translate:
  - *The animal didn't cross the street because it was too tired*
  - *The animal didn't cross the street because it was too wide*
- "it" refers to "animal" in first case, but to "street" in second case; this is hard for traditional Seq2Seq models to model

Self-Attention

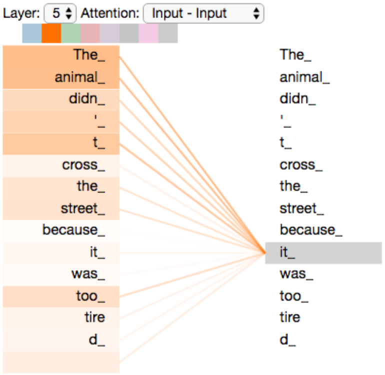

- Consider two input sentences we want to translate:
  - *The animal didn't cross the street because it was too tired*
  - *The animal didn't cross the street because it was too wide*
- "it" refers to "animal" in first case, but to "street" in second case; this is hard for traditional Seq2Seq models to model
- As the model processes each word, self-attention allows it to look at other positions in input sequence to help get a better encoding

Self-Attention

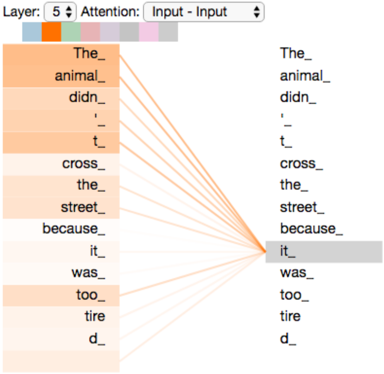

- Consider two input sentences we want to translate:
  - *The animal didn't cross the street because it was too tired*
  - *The animal didn't cross the street because it was too wide*
- "it" refers to "animal" in first case, but to "street" in second case; this is hard for traditional Seq2Seq models to model
- As the model processes each word, self-attention allows it to look at other positions in input sequence to help get a better encoding
- Recall RNNs: we now no longer need to maintain a hidden state to incorporate representation of previous words/vectors!

Self-Attention

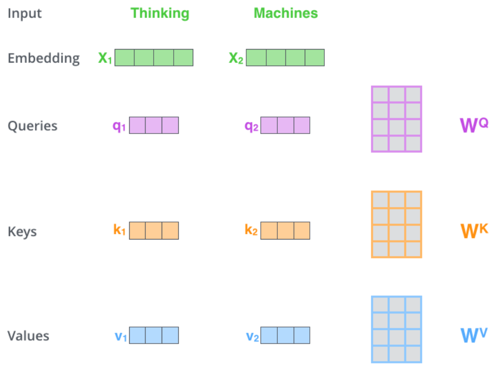

- **STEP 1:** Create three vectors from encoder's input vector (x_i):
  - Query vector (q_i)
  - Key vector (k_i)
  - Value vector (v_i)

Self-Attention

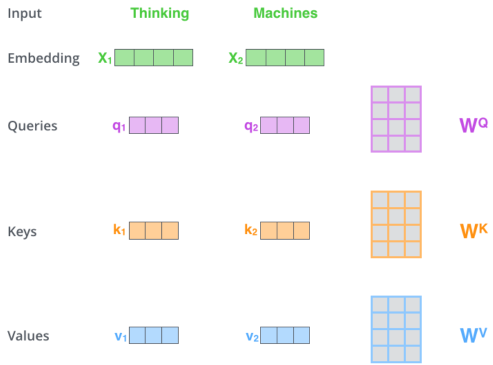

- **STEP 1:** Create three vectors from encoder's input vector (x_i):
  - Query vector (q_i)
  - Key vector (k_i)
  - Value vector (v_i)
- These are created by multiplying input with weight matrices W^Q, W^K, W^V, learned during training

Self-Attention

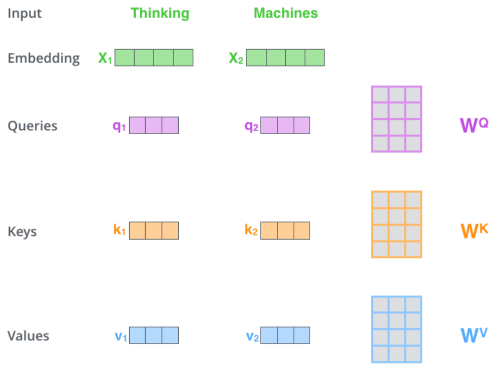

- **STEP 1:** Create three vectors from encoder's input vector (x_i):
  - Query vector (q_i)
  - Key vector (k_i)
  - Value vector (v_i)
- These are created by multiplying input with weight matrices W^Q, W^K, W^V, learned during training
- In the paper, $q, k, v \in \mathbb{R}^{64}$ and $x \in \mathbb{R}^{512}$

Self-Attention

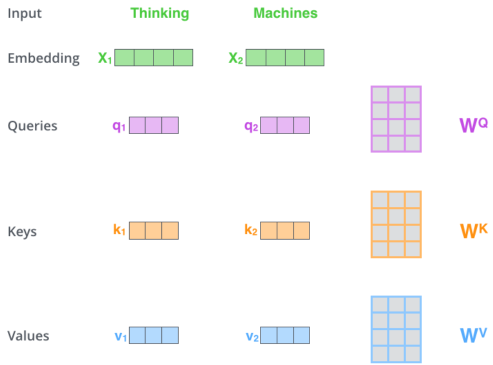

- **STEP 1:** Create three vectors from encoder's input vector (x_i):
  - Query vector (q_i)
  - Key vector (k_i)
  - Value vector (v_i)
- These are created by multiplying input with weight matrices W^Q, W^K, W^V, learned during training
- In the paper, $q, k, v \in \mathbb{R}^{64}$ and $x \in \mathbb{R}^{512}$
- Do q, k, v always have to be smaller than x?

Self-Attention

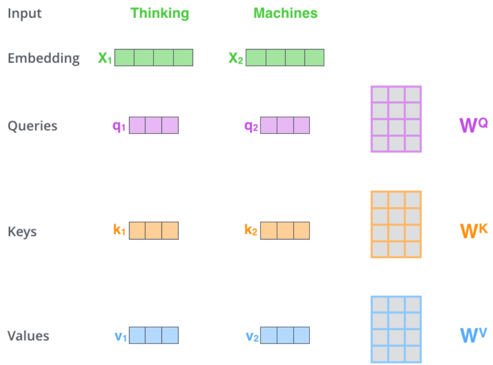

- **STEP 1:** Create three vectors from encoder's input vector (x_i):
  - Query vector (q_i)
  - Key vector (k_i)
  - Value vector (v_i)
- These are created by multiplying input with weight matrices W^Q, W^K, W^V, learned during training
- In the paper, $q, k, v \in \mathbb{R}^{64}$ and $x \in \mathbb{R}^{512}$
- Do q, k, v always have to be smaller than x? No, this was done perhaps to make computation of multi-headed attention constant
- What are the dimensions of W^Q, W^K, W^V?

Self-Attention

- **STEP 2:** Calculate self-attention scores - score all words of input sentence against themselves; how?

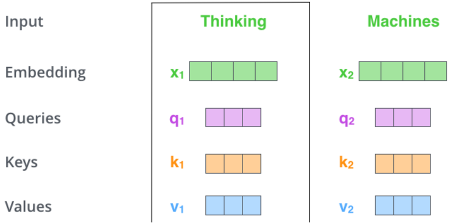

Self-Attention

- **STEP 2:** Calculate self-attention scores - score all words of input sentence against themselves; how?
- By taking dot product of **query vector** with **key vector** of respective words

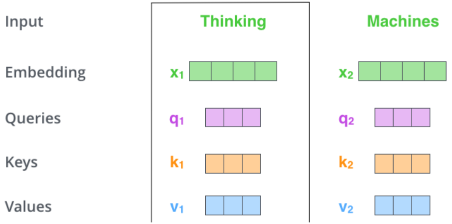

Self-Attention

- **STEP 2:** Calculate self-attention scores - score all words of input sentence against themselves; how?
- By taking dot product of **query vector** with **key vector** of respective words
- E.g. for input "Thinking", first score would be $q_1 \cdot k_1$ (with itself); second score would be dot product of $q_1 \cdot k_2$ (with "Machines"), and so on

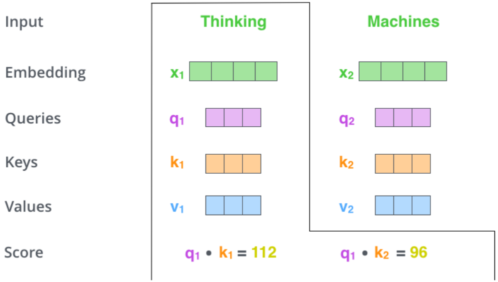

Self-Attention

- **STEP 2:** Calculate self-attention scores - score all words of input sentence against themselves; how?
- By taking dot product of **query vector** with **key vector** of respective words
- E.g. for input "Thinking", first score would be $q_1 \cdot k_1$ (with itself); second score would be dot product of $q_1 \cdot k_2$ (with "Machines"), and so on
- Scores then divided by $\sqrt{length(k)}$

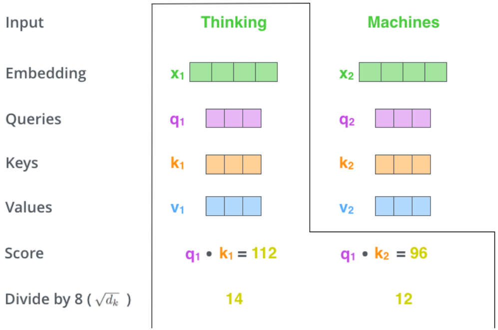

Self-Attention

- **STEP 2:** Calculate self-attention scores - score all words of input sentence against themselves; how?
- By taking dot product of **query vector** with **key vector** of respective words
- E.g. for input "Thinking", first score would be $q_1 \cdot k_1$ (with itself); second score would be dot product of $q_1 \cdot k_2$ (with "Machines"), and so on
- Scores then divided by $\sqrt{length(k)}$
- This is **Scaled Dot-Product Attention**, recall from W9P1; this design choice leads to more stable gradients

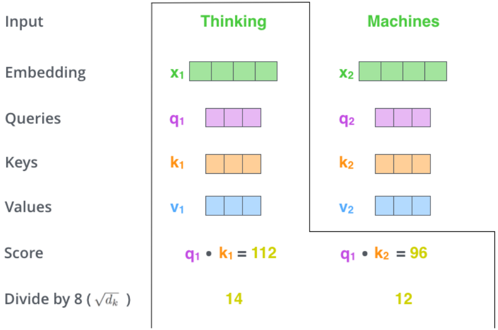

Self-Attention

- **STEP 3:** Softmax used to get normalized probability scores; determines how much each word will be expressed at this position

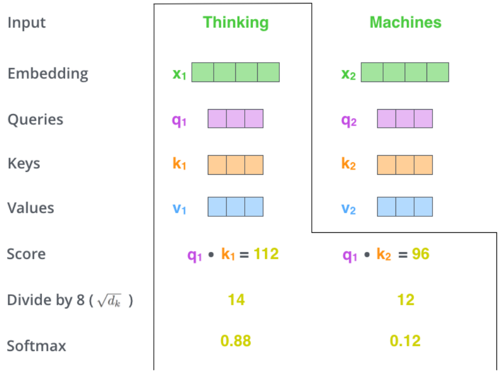

Self-Attention

- **STEP 3:** Softmax used to get normalized probability scores; determines how much each word will be expressed at this position
- Clearly, word at this position will have highest softmax score, but sometimes it's useful to attend to another word that is relevant

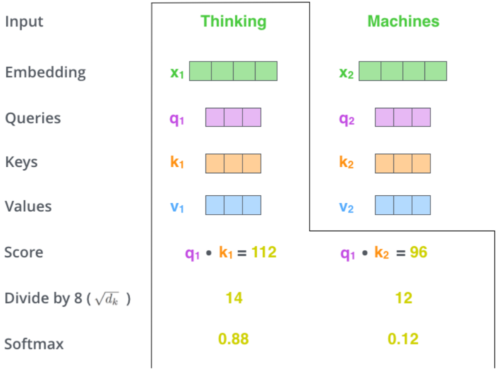

- **STEP 3:** Softmax used to get normalized probability scores; determines how much each word will be expressed at this position
- Clearly, word at this position will have highest softmax score, but sometimes it's useful to attend to another word that is relevant
- **STEP 4:** Multiply each **value vector** by softmax score; why? Keep values of word(s) we want to focus on intact, and drown out irrelevant words

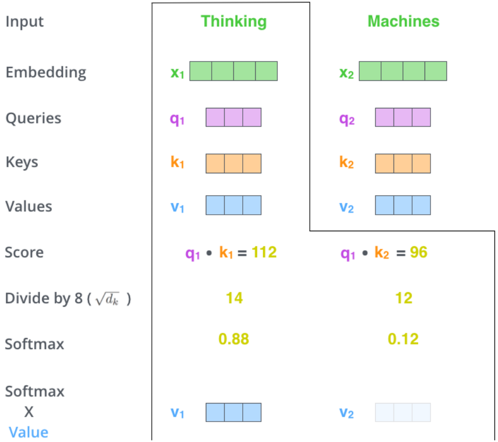

Self-Attention

- **STEP 3:** Softmax used to get normalized probability scores; determines how much each word will be expressed at this position
- Clearly, word at this position will have highest softmax score, but sometimes it's useful to attend to another word that is relevant
- **STEP 4:** Multiply each **value vector** by softmax score; why? Keep values of word(s) we want to focus on intact, and drown out irrelevant words
- **STEP 5:** Sum up weighted value vectors $\rightarrow$ produces output of self-attention layer at this position (for first word)

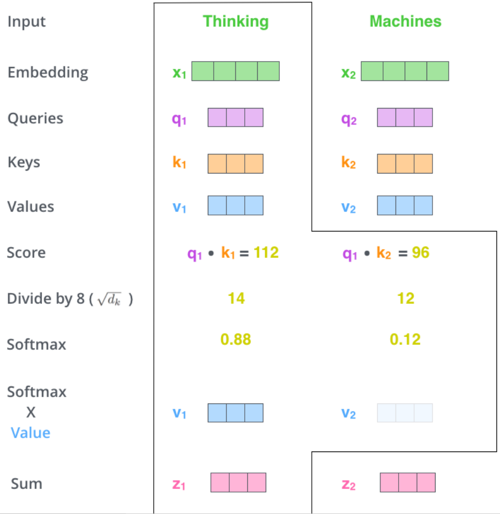

Self-Attention: Illustration

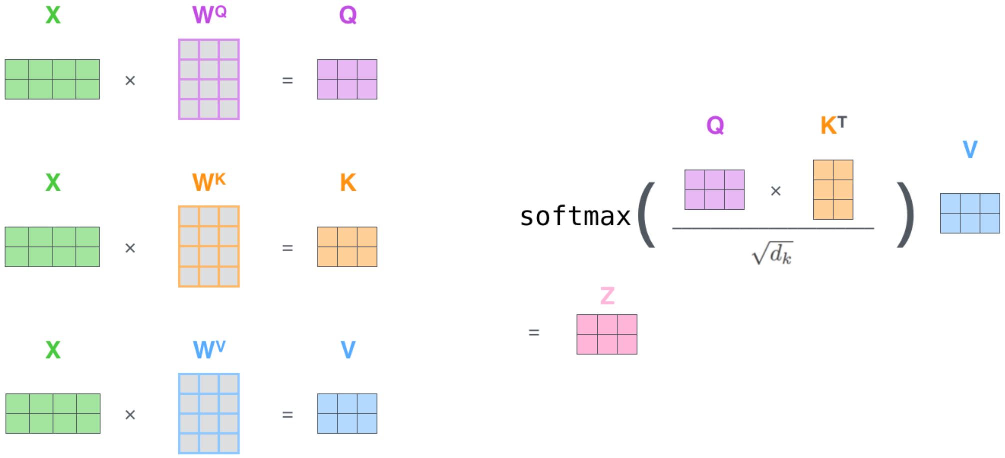

Multi-Head Attention

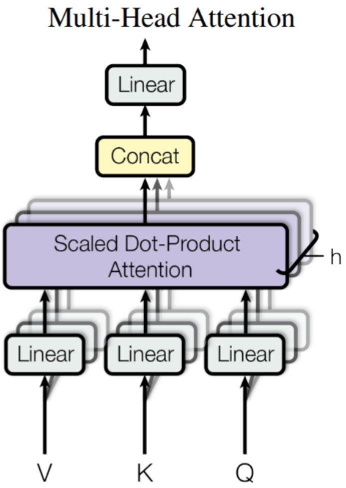

- Improves performance of the attention layer in two ways:

Multi-Head Attention

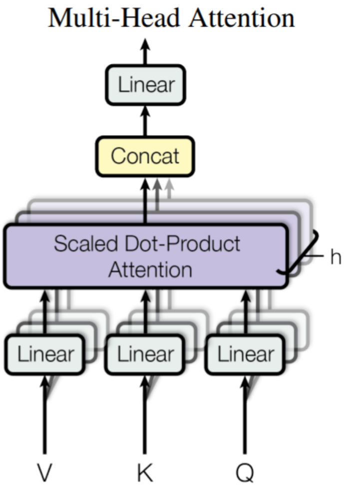

- Improves performance of the attention layer in two ways:
  - Expands model's ability to focus on different positions. In example above, z_1 contains a bit of every other encoding, but dominated by actual word itself

Multi-Head Attention

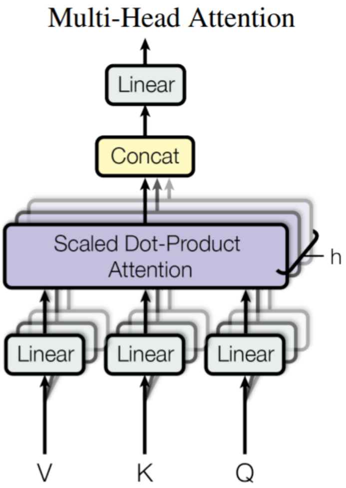

- Improves performance of the attention layer in two ways:
  - Expands model's ability to focus on different positions. In example above, z_1 contains a bit of every other encoding, but dominated by actual word itself
  - Gives attention layer multiple "*representation subspaces*"; we have not one, but multiple sets of Query/Key/Value weight matrices; after training, each set is used to project input embeddings into different representation subspaces

Credit: Vaswani et al, Attention is All You Need, NeurIPS 2017

Multi-Head Attention: Illustration

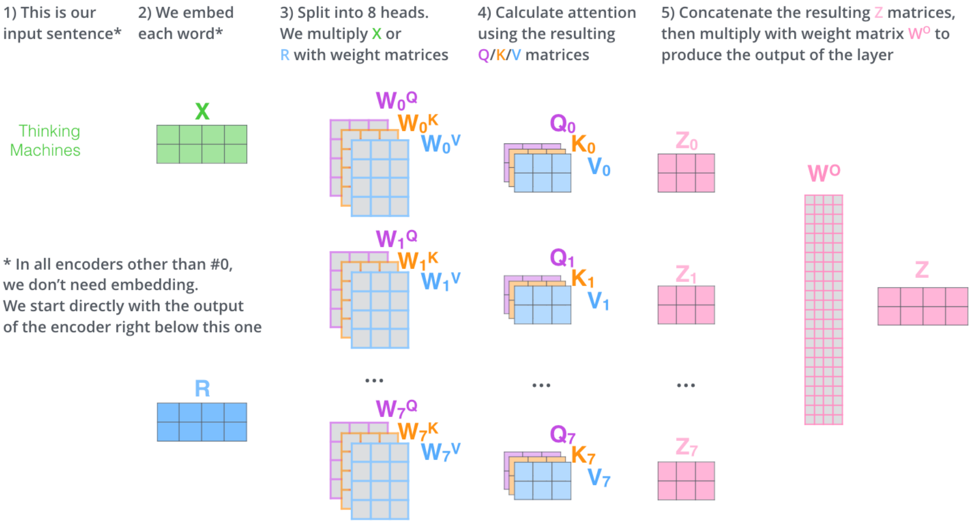

Positional Encoding

- Unlike RNN and CNN encoders, attention encoder outputs do not depend on order of inputs (Why?)

Positional Encoding

- Unlike RNN and CNN encoders, attention encoder outputs do not depend on order of inputs (Why?)
- But order of sequence conveys important information for machine translation tasks and language modeling

Positional Encoding

- Unlike RNN and CNN encoders, attention encoder outputs do not depend on order of inputs (Why?)
- But order of sequence conveys important information for machine translation tasks and language modeling
- The idea: Add positional information of input token in the sequence into input embedding vectors

$$PE_{pos,2i} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{emb}}}}\right) \qquad PE_{pos,2i+1} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{emb}}}}\right)$$

Positional Encoding

- Unlike RNN and CNN encoders, attention encoder outputs do not depend on order of inputs (Why?)
- But order of sequence conveys important information for machine translation tasks and language modeling
- The idea: Add positional information of input token in the sequence into input embedding vectors

$$PE_{pos,2i} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{emb}}}}\right) \qquad PE_{pos,2i+1} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{emb}}}}\right)$$

- Final input embeddings are concatenation of learnable embedding and positional encoding

Encoder

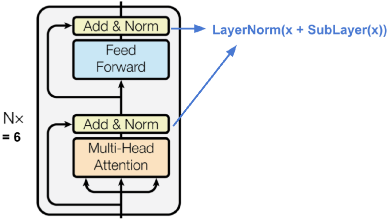

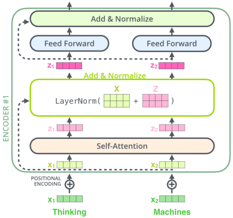

- Stack of N=6 identical layers

Credit: "Attention? Attention!" by Lilian Weng

Encoder

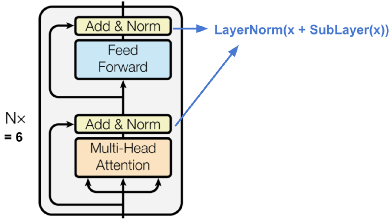

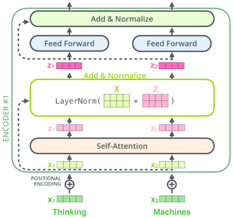

- Stack of N=6 identical layers
- Each layer has a **multi-head self-attention layer** and a simple position-wise fully connected **feedforward network**

Credit: "Attention? Attention!" by Lilian Weng

Encoder

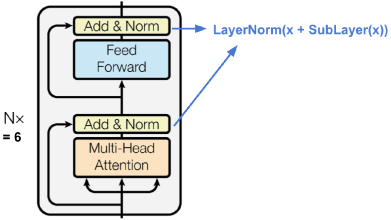

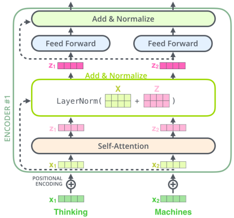

- Stack of N=6 identical layers
- Each layer has a **multi-head self-attention layer** and a simple position-wise fully connected **feedforward network**
- Each sub-layer has a **residual** connection and **layer-normalization**; all sub-layers output data of same dimension $d_{model} = 512$

Credit: "Attention? Attention!" by Lilian Weng

Decoder

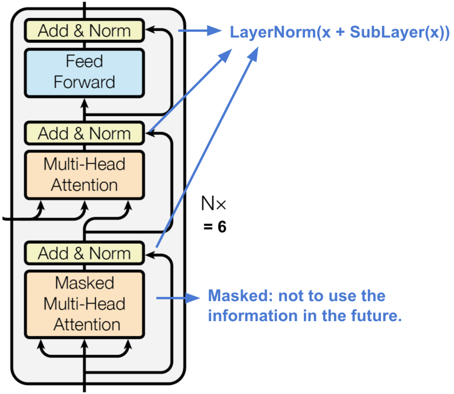

- Stack of **N=6** identical layers

Credit: "Attention? Attention!" by Lilian Weng

Decoder

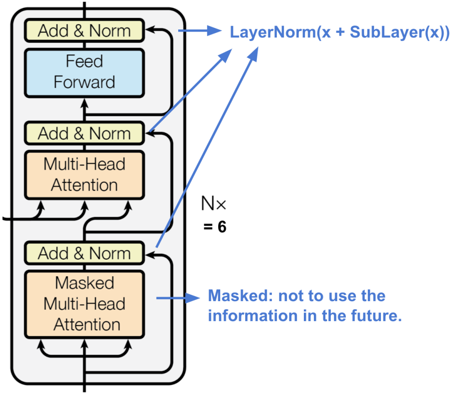

- Stack of **N=6** identical layers
- Each layer has two sub-layers of **multi-head attention** mechanisms and one sub-layer of fully-connected **feedforward network**

Credit: "Attention? Attention!" by Lilian Weng

Decoder

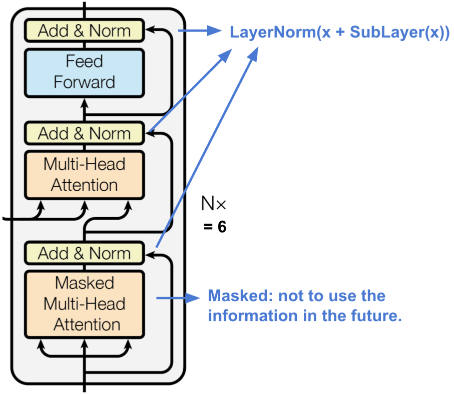

- Stack of **N=6** identical layers
- Each layer has two sub-layers of **multi-head attention** mechanisms and one sub-layer of fully-connected **feedforward network**
- Similar to encoder, each sub-layer adopts a **residual connection** and a **layer-normalization**

Credit: "Attention? Attention!" by Lilian Weng

Decoder

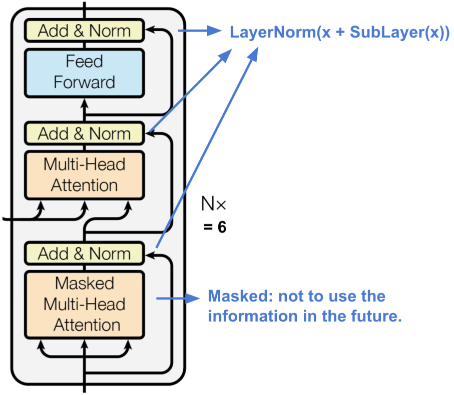

- Stack of **N=6** identical layers
- Each layer has two sub-layers of **multi-head attention** mechanisms and one sub-layer of fully-connected **feedforward network**
- Similar to encoder, each sub-layer adopts a **residual connection** and a **layer-normalization**
- First multi-head attention sub-layer is modified to prevent positions from attending to subsequent positions, as we don't want to look into future of target sequence when predicting current position

Credit: "Attention? Attention!" by Lilian Weng

Transformers: Full Architecture

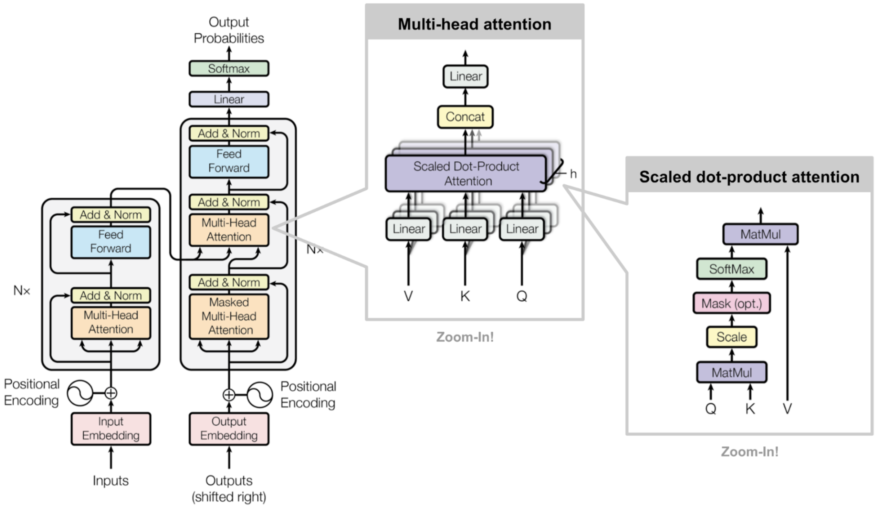

Credit: "Attention? Attention!" by Lilian Weng

Homework

Readings

- Watch the Transformers in Action video provided in the week's lecture materials
- The Illustrated Transformer article by Jay Alammar
- A detailed explanation of positional encoding by Amirhossein Kazemnejad
- For more information: Attention is All You Need paper by Vaswani, et al. (NeurIPS 2017)

Questions

- Are transformers faster or slower than LSTMs? What is the reason for your opinion?