

# Text-Conditioned and Latent Diffusion Models

Vineeth N Balasubramanian

Department of Computer Science and Engineering  
Indian Institute of Technology, Hyderabad



# Text-Conditioned Generation

- In the last lecture, we saw classifier and classifier-free guidance for class-based generation. We used an additional input  $y$  (class label) to model the conditional distribution  $p(x|y)$ . Which allows us to generate the desired image given the conditioning signal.
- We can also extend this concept to text and replace the class labels with text sequences for text-to-image generation diffusion models
- We can leverage existing language or vision-language models instead of a classifier to achieve this

# Why Text-Conditioned?

- Class-based guidance is incapable of generating complex images
- E.g. *“a dog sitting on the table”*, *“an astronaut sitting on a horse on the moon”*
- To enhance the precision of this generation task, we require greater controllability over the diffusion process → we require a text-conditioned diffusion model



“a high-quality oil painting of a psychedelic hamster dragon”



“an illustration of albert einstein wearing a superhero costume”



“a red cube on top of a blue cube”




“a stained glass window of a panda eating bamboo”

---

Image Credit: Nichol et al, *GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models*, ICML 2022

# GLIDE from OpenAI

- GLIDE from OpenAI proposed caption and CLIP guidance for the text-conditioned diffusion model, trained on hundreds of millions of paired datasets
- The model is trained by feeding a text prompt into a massive diffusion model as a condition
- $64 \times 64$  base diffusion model, which uses UNet architecture
- Token embedding from CLIP text encoder is placed into the class embedding of the UNet model



---

Nichol et al, *GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models*, ICML 2022

# Text-Conditioned Diffusion Model

## Understanding CLIP, a vision-language foundation model:

- CLIP: Contrastive Language-Image Pre-training.
- Projects given text and image pairs into the same embedding space via a text and image encoder, uses cosine similarity distance function to train model in a contrastive manner
- Can efficiently learn visual concepts in the form of text via natural language supervision

(More in the next week...)

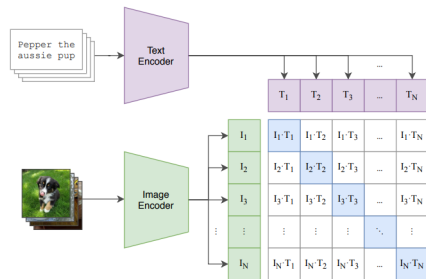


Image Credit: Radford et al, *Learning Transferable Visual Models From Natural Language Supervision*, ICML 2021

# CLIP-Guided Text-Conditioned Diffusion Model

- To use CLIP guidance, we need to re-train CLIP on the noised dataset
- Given an image  $x$  and a prompt  $c$ , a CLIP model computes the alignment via cosine similarity between  $x$  and  $c$ , indicating how similar the image and the prompt are

# CLIP-Guided Text-Conditioned Diffusion Model

- To use CLIP guidance, we need to re-train CLIP on the noised dataset
- Given an image  $x$  and a prompt  $c$ , a CLIP model computes the alignment via cosine similarity between  $x$  and  $c$ , indicating how similar the image and the prompt are
- To get the guidance signal, we use CLIP similarity score and **calculate the gradient of this score w.r.t the noised image  $x_t$** :

$$\hat{\mu}_{\theta}(x_t|c) = \mu_{\theta}(x_t|c) + s \cdot \Sigma_{\theta}(x_t|c) \nabla_{x_t}(f(x_t) \cdot g(c))$$

where:

- $f(x_t)$  - CLIP Image Encoder;  $g(c)$  - CLIP Text Encoder
- $s$  - guidance weight (same as  $\gamma$  in classifier and classifier-free guidance)
- $\Sigma_{\theta}(x_t|c)$  - covariance matrix

# CLIP-Guided Text-Conditioned Diffusion Model

- To use CLIP guidance, we need to re-train CLIP on the noised dataset
- Given an image  $x$  and a prompt  $c$ , a CLIP model computes the alignment via cosine similarity between  $x$  and  $c$ , indicating how similar the image and the prompt are
- To get the guidance signal, we use CLIP similarity score and **calculate the gradient of this score w.r.t the noised image  $x_t$** :

$$\hat{\mu}_{\theta}(x_t|c) = \mu_{\theta}(x_t|c) + s \cdot \Sigma_{\theta}(x_t|c) \nabla_{x_t}(f(x_t) \cdot g(c))$$

where:

- $f(x_t)$  - CLIP Image Encoder;  $g(c)$  - CLIP Text Encoder
- $s$  - guidance weight (same as  $\gamma$  in classifier and classifier-free guidance)
- $\Sigma_{\theta}(x_t|c)$  - covariance matrix

Note: Without re-training CLIP, performance is sub-optimal



# GLIDE: Text-Conditioned Diffusion Model

- Classifier-Free Diffusion Guidance with caption conditioning (rather than class conditioning):

$$\hat{\epsilon}_{\theta}(x_t|c) = \epsilon_{\theta}(x_t|\emptyset) + s \cdot (\epsilon_{\theta}(x_t|c) - \epsilon_{\theta}(x_t|\emptyset))$$

- CLIP-guided Diffusion Guidance:

$$\mu \epsilon_{\theta}(x_t|c) = \mu \theta(x_t|c) + s \cdot \sum_{\theta} (x_t|c) \nabla_{x_t} (f(x_t) \cdot g(c))$$

Caption conditioning worked better in their results

## Pixel-Space Diffusion Models: Limitations

- Input and output sizes are same, making network parameters count large
- Requires significant computation power and long time to train
- Slow at sampling or inference time
- Requires a large amount of GPU memory


## Pixel-Space Diffusion Models: Limitations

- Input and output sizes are same, making network parameters count large
- Requires significant computation power and long time to train
- Slow at sampling or inference time
- Requires a large amount of GPU memory

How can we reduce the computation cost and sampling time for diffusion models?

# Towards Latent Diffusion Models


- Rather than train the diffusion model in pixel space, train in latent space!
- Majority of pixels within an image represent insignificant details that may not have semantic relevance
- Latent diffusion models provide faster training and sampling, with lower memory requirements and lower cost of computation



Rombach et al, High-Resolution Image Synthesis with Latent Diffusion Models, CVPR 2022

# Understanding the VQ-VAE

- VQ-VAE uses discrete latent variables instead of continuous normal distribution in VAEs
- Encoder network takes image  $x$  and encodes into  $z_e$
- VQ layer takes  $z_e$  and samples embedding from a dictionary based on distance  $\nabla_z L$  to output  $z_q$
- Decoder network takes  $z_q$  and outputs  $x'$  to recreate input  $x$



# Understanding the VQ-VAE

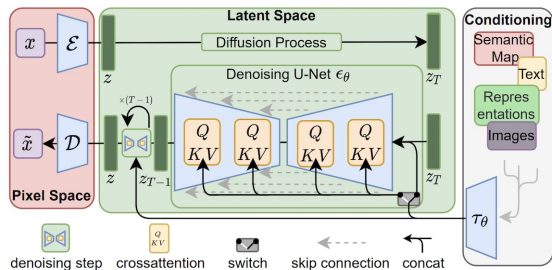
- Loss Function:

$$\mathcal{L} = \log p(x|z_q(x)) + \|\text{sg}[z_e(x)] - e\|^2 + \beta \|z_e(x) - \text{sg}[e]\|^2$$

- **Reconstruction loss:**  $\log p(x|z_q(x))$  is used to optimize the encoder and decoder networks
- **Embedding space loss:**  $\|\text{sg}[z_e(x)] - e\|^2$ , a dictionary learning algorithm uses an  $l_2$  error to move the embedding vectors  $e_i$  towards the encoder output (sg represents stop gradient)
- **Commitment loss:**  $\beta \|z_e(x) - \text{sg}[e]\|^2$  is used to control the volume of embedding space and make sure the encoder commits to an embedding
- $\beta$  is a hyperparameter that controls how much to weigh

# Latent Diffusion Models (LDMs)

- To achieve LDMs:
  - First compress image  $x$  using VQ-VAE encoder  $\mathcal{E}$  into a lower dimensional latent embedding  $z$
  - Then,  $z$  is fed into the U-Net model to learn to generate latent (i.e., compressed representations) of image which are then decoded into image  $\hat{x}$  via the VQ-VAE decoder  $D$
  - To make it a conditional model, a  $T_\theta$  encoder is used to encode (text, label or image) into embedding and pass it to cross-attention layers of UNet as a guidance signal



# Training LDMs

- **Step 1:** Train a VQ-VAE model (or use pre-trained model) on the dataset
- **Step 2:** Train  $T_\theta$  encoder (or use pre-trained model) on the dataset
  - In LDM training, both VQ-VAE Encoder and Decoder are frozen
- **Step 3:** Train Latent Diffusion UNet model and  $\tau_\theta$  encoder model using the loss function:

$$|\epsilon - \epsilon_\theta(z_t, t, \tau_\theta(y))|^2$$

- Jointly optimize both  $\tau_\theta$  and  $\epsilon_\theta$  using above loss function



# Popular Diffusion Models

- Stable Diffusion: <https://huggingface.co/spaces/stabilityai/stable-diffusion>
- DALLÉ-2: <https://github.com/lucidrains/DALLÉ2-pytorch>
- Imagen: <https://imagen.research.google/>
- SORA: <https://github.com/hpcaitech/Open-Sora>
- DreamBooth: <https://huggingface.co/docs/diffusers/en/training/dreambooth>
- ControlNet: <https://huggingface.co/docs/diffusers/en/using-diffusers/controlnet>

# Homework

## Readings

- Lilian Weng, What are Diffusion Models?
- Text-to-Image: Diffusion, Text Conditioning, Guidance, Latent Space
- (YouTube video) OpenAI GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models
- (YouTube video) VQ-VAE — Everything you need to know about it.